

Here we will predict the quality of wine on the basis of given features. We use the wine quality dataset available on Internet for free. This dataset has the fundamental features which are responsible for affecting the quality of the wine. By the use of several Machine learning models, we will predict the quality of the wine.

Our dataset that we will use consists of the following features:

1. type: whether white or red wine
2. fixed acidity: fixed acidity value
3. volatile acidity: volatile acidity value
4. citric acid: citric acid value
5. residual sugar: residual sugar value
6. chlorides: chloride value
7. free sulfur dioxide: free sulfur dioxide value
8. total sulfur dioxide: total sulfur dioxide value
9. density: density value of wine
10. pH: pH value of wine
11. sulphates: sulphates value
12. alcohol: Alcohol value
13. quality: target quality of wine which range from 0 to 10

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler

import warnings
warnings.filterwarnings("ignore")

from imblearn.over_sampling import RandomOverSampler
from imblearn.under_sampling import RandomUnderSampler
from imblearn.under_sampling import TomekLinks
from collections import Counter
from imblearn.over_sampling import SMOTE

from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score, confusion_matrix
```

## Reading data

```
In [2]: df = pd.read_csv("/Users/HP/Desktop/winequalityN.csv")
df
```

```
Out[2]:
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	white	7.0	0.270	0.36	20.7	0.045	45.0	170.0	1.00100	3.00	0.45	8.8	6
1	white	6.3	0.300	0.34	1.6	0.049	14.0	132.0	0.99400	3.30	0.49	9.5	6

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
2	white	8.1	0.280	0.40	6.9	0.050	30.0	97.0	0.99510	3.26	0.44	10.1	6
3	white	7.2	0.230	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	0.40	9.9	6
4	white	7.2	0.230	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	0.40	9.9	6
...	...	...	...	...	...	...	...	...	...	...	...	...	...
6492	red	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	5
6493	red	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	NaN	11.2	6
6494	red	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	6
6495	red	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	5
6496	red	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	6

6497 rows × 13 columns

In [3]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   type                                  6497 non-null   object
1   fixed acidity                         6487 non-null   float64
2   volatile acidity                     6489 non-null   float64
3   citric acid                          6494 non-null   float64
4   residual sugar                       6495 non-null   float64
5   chlorides                           6495 non-null   float64
6   free sulfur dioxide                  6497 non-null   float64
7   total sulfur dioxide                 6497 non-null   float64
8   density                             6497 non-null   float64
9   pH                                   6488 non-null   float64
10  sulphates                           6493 non-null   float64
11  alcohol                             6497 non-null   float64
12  quality                             6497 non-null   int64
dtypes: float64(11), int64(1), object(1)
memory usage: 660.0+ KB
```

In [4]:

```
df.describe().T
```

Out[4]:

	count	mean	std	min	25%	50%	75%	max
<b>fixed acidity</b>	6487.0	7.216579	1.296750	3.80000	6.40000	7.00000	7.70000	15.90000
<b>volatile acidity</b>	6489.0	0.339691	0.164649	0.08000	0.23000	0.29000	0.40000	1.58000
<b>citric acid</b>	6494.0	0.318722	0.145265	0.00000	0.25000	0.31000	0.39000	1.66000
<b>residual sugar</b>	6495.0	5.444326	4.758125	0.60000	1.80000	3.00000	8.10000	65.80000
<b>chlorides</b>	6495.0	0.056042	0.035036	0.00900	0.03800	0.04700	0.06500	0.61100
<b>free sulfur dioxide</b>	6497.0	30.525319	17.749400	1.00000	17.00000	29.00000	41.00000	289.00000
<b>total sulfur dioxide</b>	6497.0	115.744574	56.521855	6.00000	77.00000	118.00000	156.00000	440.00000
<b>density</b>	6497.0	0.994697	0.002999	0.98711	0.99234	0.99489	0.99699	1.03898

	count	mean	std	min	25%	50%	75%	max
<b>pH</b>	6488.0	3.218395	0.160748	2.72000	3.11000	3.21000	3.32000	4.01000
<b>sulphates</b>	6493.0	0.531215	0.148814	0.22000	0.43000	0.51000	0.60000	2.00000
<b>alcohol</b>	6497.0	10.491801	1.192712	8.00000	9.50000	10.30000	11.30000	14.90000
<b>quality</b>	6497.0	5.818378	0.873255	3.00000	5.00000	6.00000	6.00000	9.00000

## checking for nulls

```
In [5]: df.isnull().sum()
```

```
Out[5]: type                0
fixed acidity            10
volatile acidity         8
citric acid              3
residual sugar           2
chlorides                2
free sulfur dioxide      0
total sulfur dioxide     0
density                 0
pH                      9
sulphates               4
alcohol                 0
quality                 0
dtype: int64
```

- All of the features that has missing values are numerical so they can be replaced by their mean or median
- Our dataset is large so we can drop them or replace them by a suitable way
- Decided to replace them by their mean

```
In [6]: for col, value in df.items():
        if col != 'type':
            df[col] = df[col].fillna(df[col].mean())
```

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   type                  6497 non-null  object
 1   fixed acidity          6497 non-null  float64
 2   volatile acidity       6497 non-null  float64
 3   citric acid            6497 non-null  float64
 4   residual sugar         6497 non-null  float64
 5   chlorides              6497 non-null  float64
 6   free sulfur dioxide    6497 non-null  float64
 7   total sulfur dioxide   6497 non-null  float64
 8   density                6497 non-null  float64
 9   pH                    6497 non-null  float64
10   sulphates              6497 non-null  float64
11   alcohol                6497 non-null  float64
12   quality                6497 non-null  int64
dtypes: float64(11), int64(1), object(1)
memory usage: 660.0+ KB
```

# Encoding of (type)

```
In [8]: df['type'].unique()
```

Out[8]: array(['white', 'red'], dtype=object)

```
In [9]: df['type'] = [1 if i == 'white' else 0 for i in df.type]
```

```
In [10]: df['type'].unique()
```

Out[10]: array([1, 0], dtype=int64)

- replaced the type values of (white,red) with values of ( 1 for white , 0 for red)

# checking for Duplicated values

```
In [11]: df.duplicated().sum()
```

Out[11]: 1168

```
In [12]: duplicated = df.duplicated()
df[duplicated]
```

Out[12]:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
4	1	7.2	0.230	0.32	8.50	0.058	47.0	186.0	0.99560	3.19	0.40	9.9	6
5	1	8.1	0.280	0.40	6.90	0.050	30.0	97.0	0.99510	3.26	0.44	10.1	6
7	1	7.0	0.270	0.36	20.70	0.045	45.0	170.0	1.00100	3.00	0.45	8.8	6
8	1	6.3	0.300	0.34	1.60	0.049	14.0	132.0	0.99400	3.30	0.49	9.5	6
39	1	7.3	0.240	0.39	17.95	0.057	45.0	149.0	0.99990	3.21	0.36	8.6	5
...	...	...	...	...	...	...	...	...	...	...	...	...	...
6461	0	7.2	0.695	0.13	2.00	0.076	12.0	20.0	0.99546	3.29	0.54	10.1	5
6462	0	7.2	0.695	0.13	2.00	0.076	12.0	20.0	0.99546	3.29	0.54	10.1	5
6465	0	7.2	0.695	0.13	2.00	0.076	12.0	20.0	0.99546	3.29	0.54	10.1	5
6479	0	6.2	0.560	0.09	1.70	0.053	24.0	32.0	0.99402	3.54	0.60	11.3	5
6494	0	6.3	0.510	0.13	2.30	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	6

1168 rows × 13 columns

```
In [13]: len(df)
```

Out[13]: 6497

```
In [14]: df.drop_duplicates(inplace = True)
```

- They might be different wine testers for the same wine type but leaving them will affect our model results
- Decided to drop the duplicates out for better performance of the model


```
In [15]: len(df)
# 6497 - 1168 = 5329
```

```
Out[15]: 5329
```

## EDA

```
In [16]: sns.countplot(x = df['type'])
```

```
Out[16]: <Axes: xlabel='type', ylabel='count'>
```



- There is more white wine numbers than red wines


```
In [17]: def plot_histplots(dataframe):

    columns = dataframe.columns


    for column in columns:
        if column != 'type':
            plt.figure(figsize = (8,5))
            plt.hist(dataframe[column],bins = 30)
            plt.title(f'{column}')
            plt.xlabel(column)
            plt.show()
```

```
In [18]: plot_histplots(df)
```


fixed acidity




volatile acidity




citric acid




residual sugar






chlorides






free sulfur dioxide









- we can notice there is some skewness in some of the features such as 'total sulfur dioxide','free sulfur

dioxide','residual sugar'

- we need to fix the skewness to be represented in a better way.

In [19]:

```
!pip install scipy
```

```
Requirement already satisfied: scipy in c:\users\hp\appdata\roaming\python\python39\site-packages (1.10.1)
Requirement already satisfied: numpy<1.27.0,>=1.19.5 in c:\users\hp\appdata\roaming\python\python39\site-packages (from scipy) (1.24.3)
WARNING: Ignoring invalid distribution -atplotlib (c:\users\hp\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -illow (c:\users\hp\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -atplotlib (c:\users\hp\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -illow (c:\users\hp\anaconda3\lib\site-packages)
```

In [20]:

```
from scipy.stats import skew

def plot_histplots_solving_skewness(dataframe):

    columns = dataframe.columns

    for column in columns:
        if column != 'type':

            plt.figure(figsize = (8,5))


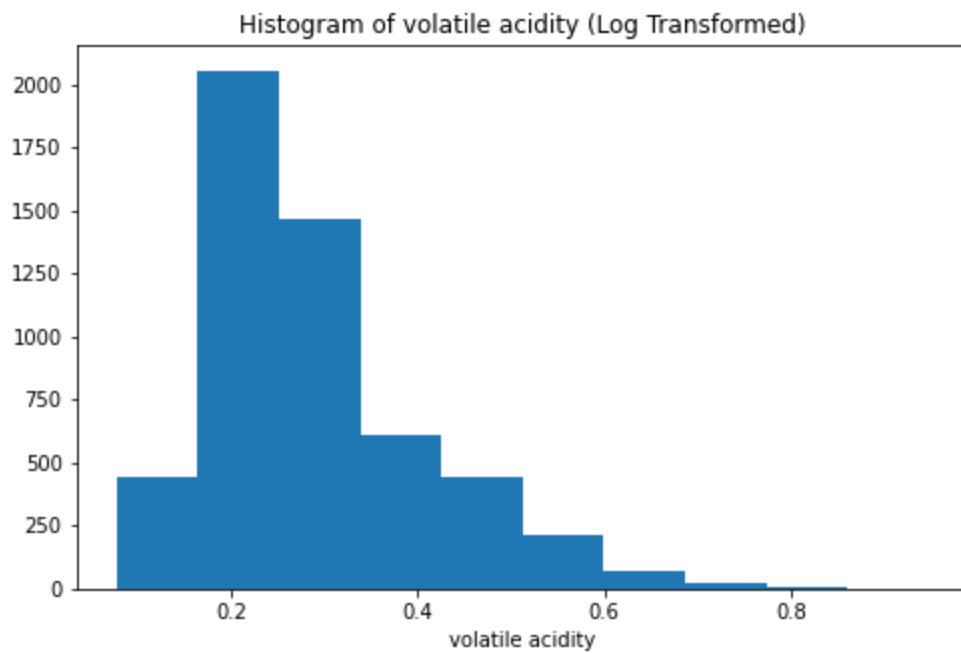
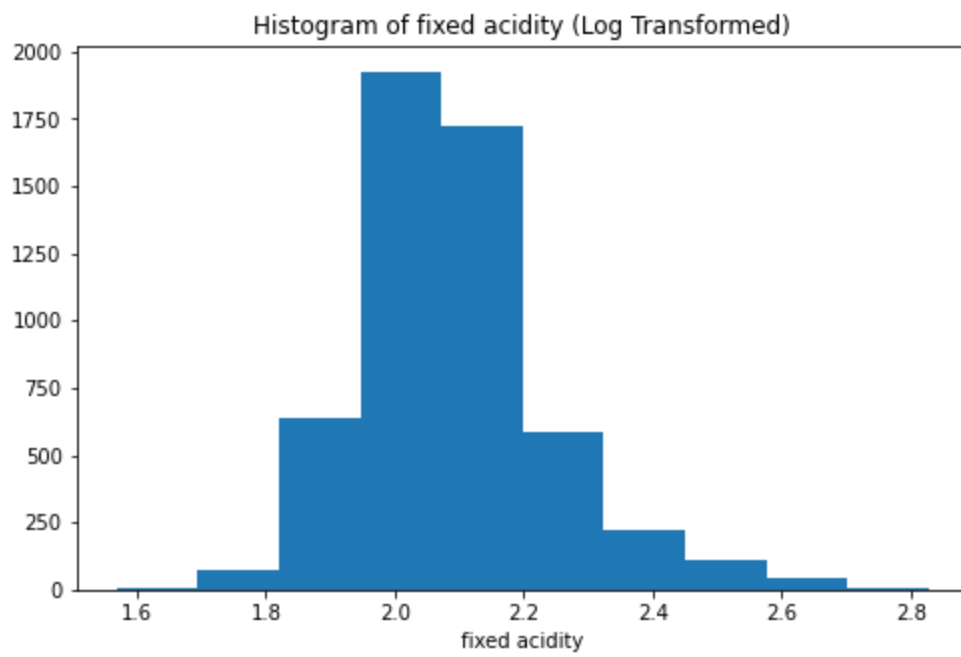
            data = dataframe[column]
            skewness = skew(data)

            if abs(skewness) > 1:
                # Apply log transformation if skewness is greater than 1
                transformed_data = np.log1p(data)
                plt.hist(transformed_data, bins=10)
                plt.title(f'Histogram of {column} (Log Transformed)')
            else:
                # Plot histogram without transformation
                plt.hist(data, bins=10)
                plt.title(f'{column}')


            plt.xlabel(column)
            plt.show()
```

In [21]:


```
plot_histplots_solving_skewness(df)
```




Histogram of residual sugar (Log Transformed)



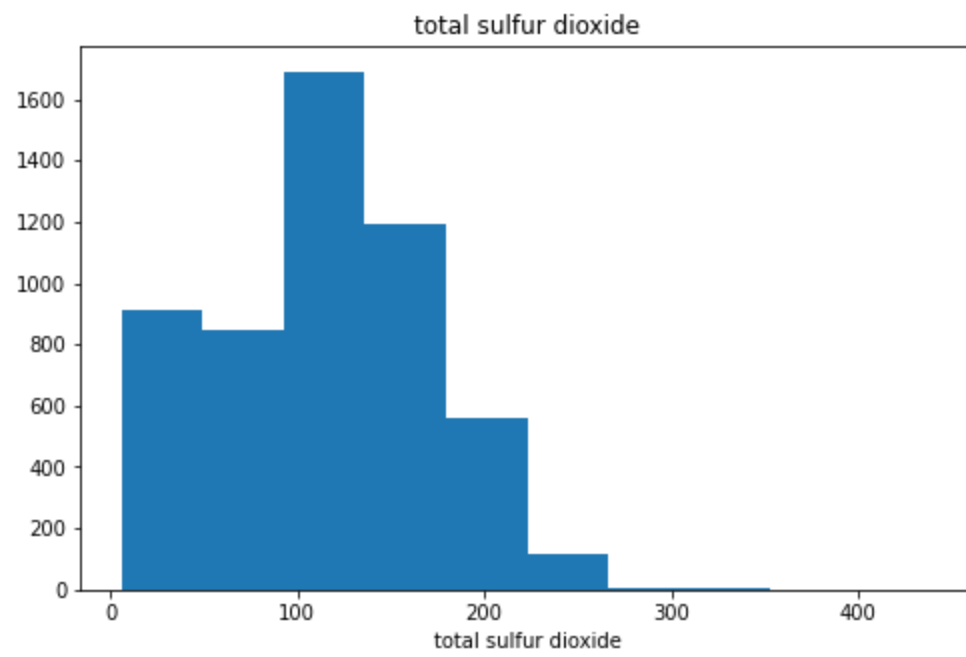


Histogram of chlorides (Log Transformed)




Histogram of free sulfur dioxide (Log Transformed)







Histogram of sulphates (Log Transformed)



alcohol



quality



- problem of skewness solved

- these are the histograms of all the features showing the distribution of each feature

# Correlation


```
In [22]: corr = df.corr()  
corr
```

Out[22]:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH
type	1.000000	-0.486281	-0.644389	0.183691	0.328820	-0.499645	0.465295	0.694181	-0.428757	-0.310338
fixed acidity	-0.486281	1.000000	0.215226	0.328758	-0.104643	0.289010	-0.281753	-0.327297	0.477858	-0.270174
volatile acidity	-0.644389	0.215226	1.000000	-0.383010	-0.164445	0.367308	-0.349388	-0.401280	0.307107	0.245559
citric acid	0.183691	0.328758	-0.383010	1.000000	0.146626	0.055081	0.132186	0.195116	0.094893	-0.343148
residual sugar	0.328820	-0.104643	-0.164445	0.146626	1.000000	-0.123254	0.399361	0.487681	0.521622	-0.233823
chlorides	-0.499645	0.289010	0.367308	0.055081	-0.123254	1.000000	-0.186836	-0.269993	0.371441	0.026176
free sulfur dioxide	0.465295	-0.281753	-0.349388	0.132186	0.399361	-0.186836	1.000000	0.720666	0.006687	-0.141344
total sulfur dioxide	0.694181	-0.327297	-0.401280	0.195116	0.487681	-0.269993	0.720666	1.000000	0.007359	-0.222514
density	-0.428757	0.477858	0.307107	0.094893	0.521622	0.371441	0.006687	0.007359	1.000000	0.034136
pH	-0.310338	-0.270174	0.245559	-0.343148	-0.233823	0.026176	-0.141344	-0.222514	0.034136	1.000000
sulphates	-0.489352	0.305803	0.226112	0.060308	-0.174795	0.404384	-0.198378	-0.274679	0.282264	0.166500
alcohol	0.057334	-0.102807	-0.065060	-0.005592	-0.306422	-0.269105	-0.170396	-0.249597	-0.668216	0.097354
quality	0.114889	-0.080554	-0.264212	0.098764	-0.057253	-0.202312	0.054456	-0.050387	-0.326978	0.039946

```
In [23]: plt.figure(figsize = (10,5))  
sns.heatmap(corr , annot = True, cmap = 'winter')
```

Out[23]: <Axes: >



- We can find that there is no features that is heavily correlated to each other
- This is great for our performance of the model
- No feature selection is needed


## Splitting of data

```
In [24]: x = df.drop(['type', 'quality'], axis = 1)
         y = df['quality']
```

## checking of imbalanced data

```
In [25]: sns.countplot(x = df['quality'])
```

```
Out[25]: <Axes: xlabel='quality', ylabel='count'>
```





- We find that the classes number of the quality diifer from one class to another by a huge difference
- we can solve the problem by applying SMOTE or RandomOverSampler.

# Handling unbalanced data

In [26]: `y.value_counts()`

Out[26]:

6	2327
5	1755
7	857
4	206
8	149
3	30
9	5

Name: quality, dtype: int64

In [27]:

```
smote = SMOTE(k_neighbors = 4)
x,y = smote.fit_resample(x,y)
```

In [28]: `y.value_counts()`

Out[28]:

6	2327
5	2327
7	2327
8	2327
4	2327
3	2327
9	2327

Name: quality, dtype: int64

- Now each class is balanced so we can move on to the train test split

## train\_test split

In [29]:

```
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size = 0.2, random_state = 42)
x_train,x_val,y_train,y_val = train_test_split(x_train,y_train,test_size = 0.2, random_state = 42)
```

## Scaling

In [30]:

```
scaler = StandardScaler()

x_train = scaler.fit_transform(x_train)
x_val = scaler.transform(x_val)
x_test = scaler.transform(x_test)
```

- We just scaled the the input features
- We applied a fit\_transform for the x\_train
- For the rest, we just applied a transform function for them (x\_val, x\_test)

## KNN

```
In [31]: # Use grid search to find best value
from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier()

params_grid = {
    'n_neighbors': [3, 4, 5, 6, 7, 8, 9, 10]
}

grid = GridSearchCV(
    knn,
    params_grid,
    cv = 5
)
grid.fit(x_train, y_train)

print(f'the best value of k = {grid.best_params_}')
```

the best value of k = {'n\_neighbors': 3}

```
In [32]: ## check overfitting

knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(x_train, y_train)

x_train_pred = knn.predict(x_train)
train_score = accuracy_score(y_train, x_train_pred)
print(f'the train score is ={train_score}')
```


x\_val\_pred = knn.predict(x\_val)  
val\_score = accuracy\_score(y\_val, x\_val\_pred)  
print(f'the valid score is ={val\_score}')

# scores are almost equal so no overfitting

the train score is =0.8839217191097467  
the valid score is =0.7736862293824319

```
In [33]: y_pred = knn.predict(x_test)
cnf_mat = confusion_matrix(y_test, y_pred)
plt.figure(figsize = (10, 5))
sns.heatmap(cnf_mat, annot = True, cmap = 'Blues')
```

Out[33]: <Axes: >



In [34]:

```
acc_score_knn = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred, average = 'macro')
precision = precision_score(y_test, y_pred, average = 'macro')
f1 = f1_score(y_test, y_pred, average = 'macro')

print(f'accuracy = {acc_score_knn * 100} %')
print(f'recall = {recall * 100} %')
print(f'precision = {precision * 100} %')
print(f'f1 score = {f1 * 100} %')
```

```
accuracy = 77.93124616329035 %
recall = 78.05802932415563 %
precision = 76.01632097959961 %
f1 score = 76.35984956920231 %
```

## Decision Tree

In [35]:

```
# Use grid search to find best value
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree

dt = DecisionTreeClassifier()

params_grid = {
    'max_depth': [3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15],
    'criterion': ['gini', 'entropy', 'log_loss']
}

grid = GridSearchCV(
    dt,
    params_grid,
    cv = 5
)
grid.fit(x_train, y_train)

print(f'the best value is = {grid.best_params_}')
```

```
the best value is = {'criterion': 'entropy', 'max_depth': 15}
```

In [36]:

```
## check overfitting

dt = DecisionTreeClassifier(criterion = 'entropy', max_depth = 11)
```

```
dt.fit(x_train,y_train)
```

```
x_train_pred = dt.predict(x_train)
train_score = accuracy_score(y_train,x_train_pred)
print(f'the train score is ={train_score}')
```

```
x_val_pred = dt.predict(x_val)
val_score = accuracy_score(y_val,x_val_pred)
print(f'the valid score is ={val_score}')
```

```
# scores are almost equal so no overfitting
```

the train score is =0.8123561013046815


the valid score is =0.6797084771768316

In [37]:

```
y_pred = dt.predict(x_test)
cnf_mat = confusion_matrix(y_test,y_pred)
plt.figure(figsize = (10,5))
sns.heatmap(cnf_mat ,annot = True, cmap = 'Blues')
y_pred
```

Out[37]:

array([4, 9, 9, ..., 4, 5, 5], dtype=int64)



In [38]:

```
acc_score_dt = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred, average = 'macro')
precision = precision_score(y_test, y_pred, average = 'macro')
f1 = f1_score(y_test,y_pred, average = 'macro')

print(f'accuracy = {acc_score_dt * 100} %')
print(f'recall = {recall * 100} %')
print(f'precision = {precision * 100} %')
print(f'f1 score = {f1 * 100} %')
```

accuracy = 67.00429711479435 %

recall = 67.16885787900334 %

precision = 65.78617560171931 %

f1 score = 66.16385891157778 %

## Random Forest

In [39]:

```
# Use grid search to find best value
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.tree import plot_tree
```

```
rf = RandomForestClassifier()
```

```
params_grid = {  
    'max_depth': [3,4,5,6,7,8,9,10],  
    'criterion': ['gini', 'entropy', 'log_loss']  
}
```

```
grid = GridSearchCV(  
    rf,  
    params_grid,  
    cv = 5  
)  
grid.fit(x_train,y_train)
```

```
print(f'the best value is = {grid.best_params_}')
```

```
the best value is = {'criterion': 'log_loss', 'max_depth': 10}
```

In [40]:

```
## check overfitting
```

```
rf = RandomForestClassifier(criterion = 'log_loss', max_depth = 10)  
rf.fit(x_train,y_train)
```

```
x_train_pred = rf.predict(x_train)  
train_score = accuracy_score(y_train,x_train_pred)  
print(f'the train score is ={train_score}')
```

```
x_val_pred = rf.predict(x_val)  
val_score = accuracy_score(y_val,x_val_pred)  
print(f'the valid score is ={val_score}')
```

```
# scores are almost near and no big difference so no overfitting
```

```
the train score is =0.9136607828089025
```


```
the valid score is =0.7733026467203682
```

In [41]:

```
y_pred = rf.predict(x_test)  
cnf_mat = confusion_matrix(y_test,y_pred)  
plt.figure(figsize = (10,5))  
sns.heatmap(cnf_mat ,annot = True, cmap = 'Blues')  
y_pred
```

Out[41]:

```
array([4, 9, 9, ..., 3, 5, 5], dtype=int64)
```



In [42]:

```
acc_score_rf = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred, average = 'macro')
precision = precision_score(y_test, y_pred, average = 'macro')
f1 = f1_score(y_test, y_pred, average = 'macro')

print(f'accuracy = {acc_score_rf * 100} %')
print(f'recall = {recall * 100} %')
print(f'precision = {precision * 100} %')
print(f'f1 score = {f1 * 100} %')
```

```
accuracy = 77.68569674647023 %
recall = 77.81073905402492 %
precision = 76.49565539041333 %
f1 score = 76.92830906875642 %
```

## Visualization of Decision tree

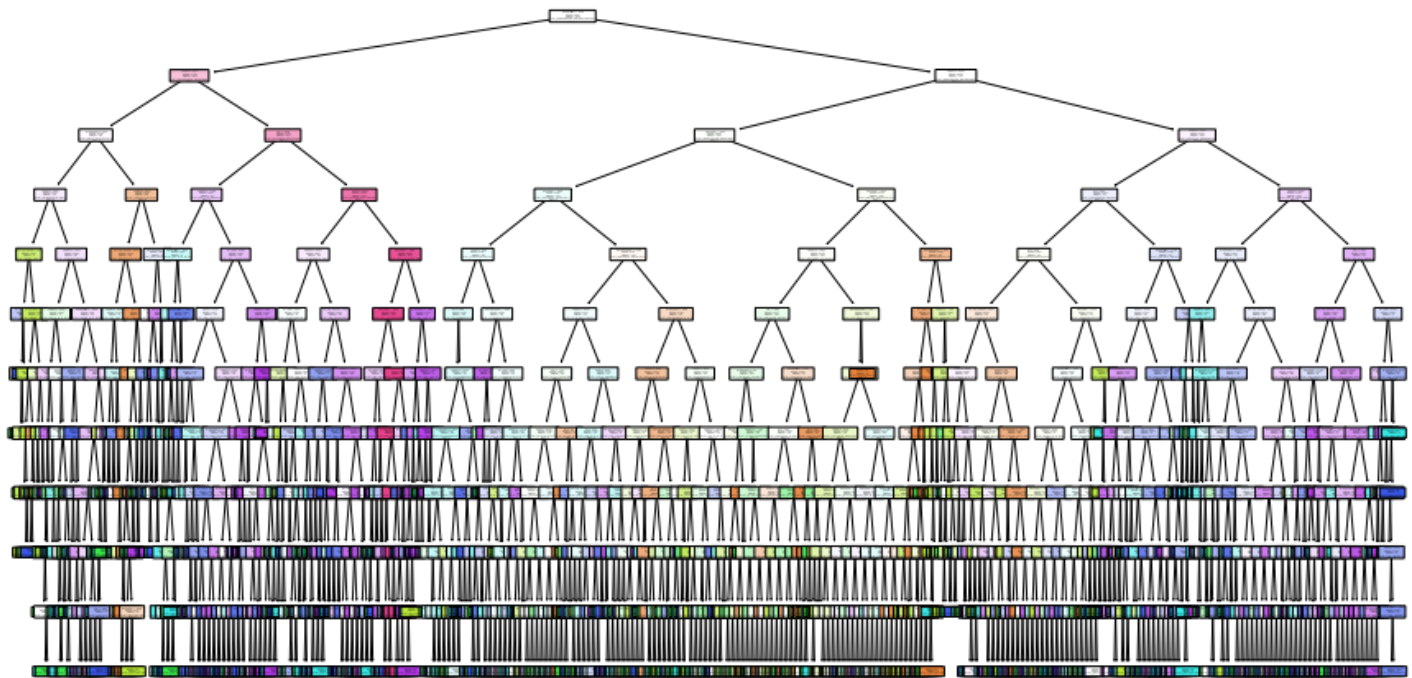
In [43]:

```
features_names = df.columns[:11]
target_names = df.columns[11]

plt.figure(figsize= (15,8))

plot_tree(
    dt,
    feature_names = features_names,
    class_names = target_names,
    filled = True, #to make it coloured
    rounded = True # to make edges rounded
)

plt.show()
```




- Here is a visualization of the decision tree in which our model produced

## Comparison of Models Evaluation (Accuracy)

```
In [44]: models_names = ['dt', 'knn', 'rf']
models_scores = [acc_score_dt, acc_score_knn, acc_score_rf]

sns.barplot(x = models_names, y = models_scores, data = df, palette = 'husl')
```

Out[44]: <Axes: >



- We can conclude from the plot above that KNN classifier scored the highest accuracy
- random forest classifier was much near to the accuracy scored by the KNN classifier than the accuracy scored by decision tree classifier.
- Decision tree classifier had the lowest score among all of the classifiers

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:



In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

